



Docker Certified Associate (DCA)

Create by: **Hadi Tayanloo**

Phone : **+98-912-8387233**

Linkedin: linkedin.com/in/htayanloo/

Session 9

- Kubernetes Management
- Kubernetes Pods
- Kubernetes Resource
- Kubernetes Volume

Kubernetes Management : Generate New token

- root@k8s-master:~|⇒ kubeadm token generate
- 9vkzol.sy0g1w465nthcviv
- root@k8s-master:~|⇒ kubeadm token create 9vkzol.sy0g1w465nthcviv --print-join-command --ttl=1h
- kubeadm join 192.168.1.200:6443 --token 9vkzol.sy0g1w465nthcviv --discovery-token-ca-cert-hash sha256:0424adc00812a3eb6c51da613c0dad4911b92be6ef7bcc0285556f7f59d83304

- root@k8s-master:~|⇒ kubeadm token list

TOKEN	TTL	EXPIRES	USAGES	DESCRIPTION	EXTRA GROUPS
9vkzol.sy0g1w465nthcviv	59m	2019-01-21T22:16:24+03:30	authentication,signing	<none>	

```
root@k8s-master:~|⇒ kubeadm token generate
9vkzol.sy0g1w465nthcviv
root@k8s-master:~|⇒ kubeadm token create 9vkzol.sy0g1w465nthcviv --print-join-command --ttl=1h
kubeadm join 192.168.1.200:6443 --token 9vkzol.sy0g1w465nthcviv --discovery-token-ca-cert-hash sha256:0424adc00812a3eb6c51da613c0dad4911b92be6ef7bcc0285556f7f59d83304
root@k8s-master:~|⇒ kubeadm token list
```

TOKEN	TTL	EXPIRES	USAGES	DESCRIPTION	EXTRA GROUPS
9vkzol.sy0g1w465nthcviv	59m	2019-01-21T22:16:24+03:30	authentication,signing	<none>	system:bootstrappers:kubeadm:default-node-token

The **conditions** field describes the status of all Running nodes.

- **OutOfDisk**
 - **True** if there is insufficient free space on the node for adding new pods, otherwise **False**
- **Ready**
 - **True** if the node is healthy and ready to accept pods, **False** if the node is not healthy and is not accepting pods, and Unknown if the node controller has not heard from the node in the last node-monitor-grace-period (default is 40 seconds)
- **MemoryPressure**
 - **True** if pressure exists on the node memory – that is, if the node memory is low; otherwise **False**
- **PIDPressure**
 - **True** if pressure exists on the processes – that is, if there are too many processes on the node; otherwise **False**
- **DiskPressure**
 - **True** if pressure exists on the disk size – that is, if the disk capacity is low; otherwise **False**
- **NetworkUnavailable**
 - **True** if the network for the node is not correctly configured, otherwise **False**
-

The **conditions** field describes the **status** of all **Running** nodes.

```
htayanloo@Linux-Home:/home/htayanloo/Downloads $ kubectl describe node
Name:           k8s-master
Roles:          master
Labels:         beta.kubernetes.io/arch=amd64
                beta.kubernetes.io/os=linux
                kubernetes.io/hostname=k8s-master
                node-role.kubernetes.io/master=
Annotations:   flannel.alpha.coreos.com/backend-data: {"VtepMAC":"8e:e5:b4:a7:90:41"}
                flannel.alpha.coreos.com/backend-type: vxlan
                flannel.alpha.coreos.com/kube-subnet-manager: true
                flannel.alpha.coreos.com/public-ip: 192.168.1.200
                kubeadm.alpha.kubernetes.io/cri-socket: /var/run/dockershim.sock
                node.alpha.kubernetes.io/ttl: 0
                volumes.kubernetes.io/controller-managed-attach-detach: true
CreationTimestamp: Mon, 21 Jan 2019 02:20:13 +0330
Taints:          node-role.kubernetes.io/master:NoSchedule
Unschedulable:  false
Conditions:    Type        Status  LastHeartbeatTime                 LastTransitionTime               Reason                               Message
                ----        -----  -----                         -----                         -----
MemoryPressure  False     Tue, 22 Jan 2019 16:20:46 +0330   Mon, 21 Jan 2019 02:20:04 +0330   KubeletHasSufficientMemory   kubelet has sufficient memory available
DiskPressure    False     Tue, 22 Jan 2019 16:20:46 +0330   Mon, 21 Jan 2019 02:20:04 +0330   KubeletHasNoDiskPressure    kubelet has no disk pressure
PIDPressure    False     Tue, 22 Jan 2019 16:20:46 +0330   Mon, 21 Jan 2019 02:20:04 +0330   KubeletHasSufficientPID     kubelet has sufficient PID available
Ready          True      Tue, 22 Jan 2019 16:20:46 +0330   Mon, 21 Jan 2019 02:24:34 +0330   KubeletReady                  kubelet is posting ready status
Addresses:    InternalIP: 192.168.1.200
              Hostname:   k8s-master
Capacity:     cpu:           2
              ephemeral-storage: 38770180Ki
              hugepages-2Mi:    0
              memory:          1882148Ki
              pods:            110
Allocatable:  cpu:           2
              ephemeral-storage: 35730597829
              hugepages-2Mi:    0
              memory:          1779748Ki
              pods:            110
```

Kubernetes node list

- root@k8s-master:~|⇒ kubectl get nodes -o wide
- | NAME | STATUS | ROLES | AGE | VERSION | INTERNAL-IP | EXTERNAL-IP | OS-IMAGE | KERNEL-VERSION | CONTAINER-RUNTIME |
|------------|--------|--------|-----|---------|---------------|-------------|-----------------------|---------------------------|-------------------|
| k8s-master | Ready | master | 19h | v1.13.2 | 192.168.1.200 | <none> | CentOS Linux 7 (Core) | 3.10.0-957.1.3.el7.x86_64 | docker://18.9.1 |
| node01 | Ready | <none> | 18h | v1.13.2 | 192.168.1.201 | <none> | CentOS Linux 7 (Core) | 3.10.0-957.1.3.el7.x86_64 | docker://18.9.1 |
| • | | | | | | | | | |
| • | | | | | | | | | |

```
root@k8s-master:~|⇒ kubectl get nodes -o wide
NAME      STATUS    ROLES      AGE     VERSION   INTERNAL-IP      EXTERNAL-IP    OS-IMAGE      KERNEL-VERSION      CONTAINER-RUNTIME
k8s-master Ready     master     19h    v1.13.2   192.168.1.200    <none>        CentOS Linux 7 (Core) 3.10.0-957.1.3.el7.x86_64  docker://18.9.1
node01    Ready     <none>    18h    v1.13.2   192.168.1.201    <none>        CentOS Linux 7 (Core) 3.10.0-957.1.3.el7.x86_64  docker://18.9.1
root@k8s-master:~|⇒
```

Kubernetes Cluster info

```
root@k8s-master:~|> kubectl cluster-info
Kubernetes master is running at https://192.168.1.200:6443
Heapster is running at https://192.168.1.200:6443/api/v1/namespaces/kube-system/services/heapster/proxy
KubeDNS is running at https://192.168.1.200:6443/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy
monitoring-influxdb is running at https://192.168.1.200:6443/api/v1/namespaces/kube-system/services/monitoring-influxdb/proxy

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
root@k8s-master:~|>
```

Get list of all Pods from any name space

```
root@k8s-master:~|⇒ kubectl get pods --all-namespaces
NAME                                READY   STATUS        RESTARTS   AGE
default    command-demo               0/1    ImagePullBackOff  0          5h46m
default    nginx-5c7588df-4hpwn      1/1    Running       0          18h
default    nginx-5c7588df-ttxsp      1/1    Running       0          9h
kube-system coredns-86c58d9df4-4sfds 1/1    Running       0          19h
kube-system coredns-86c58d9df4-qrphw 1/1    Running       0          19h
kube-system etcd-k8s-master         1/1    Running       0          19h
kube-system heapster-f64999bc-s7kx8  1/1    Running       0          8m36s
kube-system kube-apiserver-k8s-master 1/1    Running       0          19h
kube-system kube-controller-manager-k8s-master 1/1    Running       0          19h
kube-system kube-flannel-ds-amd64-lfbt5     1/1    Running       0          19h
kube-system kube-flannel-ds-amd64-vdpq5     1/1    Running       2          18h
kube-system kube-proxy-bknvk           1/1    Running       0          19h
kube-system kube-proxy-jfsvc          1/1    Running       0          18h
kube-system kube-scheduler-k8s-master   1/1    Running       0          19h
kube-system kubernetes-dashboard-57df4db6b-xblmc 0/1    CrashLoopBackOff 180         15h
kube-system kubernetes-dashboard-head-57b9585588-z49gb 0/1    ImagePullBackOff 107         15h
kube-system monitoring-influxdb-8b7d57f5c-rd9qp     1/1    Running       0          8m28s
root@k8s-master:~|⇒
```

kubectl get events

```
htayanloo@Linux-Home:/home/htayanloo/Downloads $ kubectl get events
LAST SEEN    TYPE      REASON    KIND      MESSAGE
<invalid>  Normal    Pulling   Pod       pulling image "debian"
<invalid>  Normal    BackOff   Pod       Back-off pulling image "debian"
<invalid>  Warning   Failed    Pod       Error: ImagePullBackOff
<invalid>  Warning   Failed    Pod       (combined from similar events): Failed to pull image "debian": rpc error: code = Unknown
ha256/a0/a0bd3e1c8f9eb8ff9d65828e8062ae9284b60cb83abe59fe46c74d77d88eb952/data?verify=1548650796-U6dW4p5UGaue1jAKl7uyQtThz3w%3D:
htayanloo@Linux-Home:/home/htayanloo/Downloads $
```

kubectl get events --namespace=my-namespace

```
htayanloo@Linux-Home:/home/htayanloo/Downloads $ kubectl get events --namespace=kube-system
LAST SEEN    TYPE      REASON    KIND      MESSAGE
9m46s        Normal    BackOff   Pod       Back-off pulling image "k8s.gcr.io/kubernetes-dashboard-amd64:v1.10.1"
4m49s        Warning   Failed    Pod       Error: ImagePullBackOff
<invalid>    Normal    BackOff   Pod       Back-off pulling image "kubernetesdashboarddev/kubernetes-dashboard-amd64:head"
<invalid>    Warning   Failed    Pod       Error: ImagePullBackOff
<invalid>    Warning   Failed    Pod       (combined from similar events): Error: ImagePullBackOff
htayanloo@Linux-Home:/home/htayanloo/Downloads $ █
```

kubectl get pod POD_NAME -o yaml

```
htayanloo@Linux-Home:/home/htayanloo/Downloads $ kubectl get pod nginx-5c7588df-4hpwn -o yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: "2019-01-21T00:14:32Z"
  generateName: nginx-5c7588df-
  labels:
    app: nginx
    pod-template-hash: 5c7588df
  name: nginx-5c7588df-4hpwn
  namespace: default
  ownerReferences:
  - apiVersion: apps/v1
    blockOwnerDeletion: true
    controller: true
    kind: ReplicaSet
    name: nginx-5c7588df
    uid: 86857f5f-1d11-11e9-bb57-080027808554
  resourceVersion: "9192"
  selfLink: /api/v1/namespaces/default/pods/nginx-5c7588df-4hpwn
  uid: 8689bdbf-1d11-11e9-bb57-080027808554
spec:
  containers:
  - image: nginx
    imagePullPolicy: Always
    name: nginx
    resources: {}
    terminationMessagePath: /dev/termination-log
    terminationMessagePolicy: File
    volumeMounts:
    - mountPath: /var/run/secrets/kubernetes.io/serviceaccount
      name: default-token-zcz2x
      readOnly: true
    dnsPolicy: ClusterFirst
    enableServiceLinks: true
    nodeName: node01
    priority: 0
    restartPolicy: Always
    schedulerName: default-scheduler
    securityContext: {}
    serviceAccount: default
```

POD and Container

Getting a shell to a Container

```
kubectl create -f https://k8s.io/examples/application/shell-demo.yaml
```

```
kubectl get pod shell-demo
```

```
kubectl exec -it shell-demo -- /bin/bash
```

```
root@shell-demo:/# ls /
```

Running individual commands in a Container

```
kubectl exec shell-demo env
```

```
kubectl exec shell-demo ps aux
```

```
kubectl exec shell-demo ls /
```

```
kubectl exec shell-demo cat /proc/1/mounts
```

Opening a shell when a Pod has more than one Container

If a Pod has more than one Container, use `--container` or `-c` to specify a Container in the `kubectl exec` command. For example, suppose you have a Pod named `my-pod`, and the Pod has two containers named `main-app` and `helper-app`. The following command would open a shell to the `main-app` Container.

```
kubectl exec -it my-pod --container main-app -- /bin/bash
```

Assign Memory Resources to Containers and Pods

Assign Memory Resources to Containers and Pods

Senario-1

1. kubectl create namespace mem-example
2. kubectl create -f <https://k8s.io/examples/pods/resource/memory-request-limit.yaml> --namespace=mem-example
3. kubectl get pod memory-demo --namespace=mem-example
4. kubectl get pod memory-demo --output=yaml --namespace=mem-example
5. kubectl top pod memory-demo --namespace=mem-example
6. kubectl delete pod memory-demo --namespace=mem-example

Exceed a Container's memory limit

Senario-2

1. kubectl create -f https://k8s.io/examples/pods/resource/memory-request-limit-2.yaml --namespace=mem-example
2. kubectl get pod memory-demo-2 --namespace=mem-example
3. kubectl get pod memory-demo-2 --output=yaml --namespace=mem-example
4. kubectl get pod memory-demo-2 --namespace=mem-example
5. kubectl describe pod memory-demo-2 --namespace=mem-example
6. kubectl describe nodes
7. kubectl delete pod memory-demo-2 --namespace=mem-example

Specify a memory request that is too big for your Nodes

Senario-3

1. kubectl create -f <https://k8s.io/examples/pods/resource/memory-request-limit-3.yaml> –namespace=mem-example
2. kubectl get pod memory-demo-3 –namespace=mem-example
3. kubectl describe pod memory-demo-3 –namespace=mem-example
4. kubectl delete pod memory-demo-3 --namespace=mem-example

CleanUp Scenario

- kubectl delete namespace mem-example
-



Assign CPU Resources to Containers and Pods

1.kubectl create namespace cpu-example

Assign CPU Resources to Containers and Pods

Senario-4

1. `kubectl create -f https://k8s.io/examples/pods/resource/cpu-request-limit.yaml --namespace(cpu-example)`
2. `kubectl get pod cpu-demo --namespace(cpu-example)`
3. `kubectl get pod cpu-demo --output=yaml --namespace(cpu-example)`
4. `kubectl top pod cpu-demo –namespace(cpu-example)`
5. `kubectl delete pod cpu-demo --namespace(cpu-example)`
- 6.

Specify a CPU request that is too big for your Nodes

Senario-5

1. kubectl create -f <https://k8s.io/examples/pods/resource/cpu-request-limit-2.yaml> --namespace(cpu-example)
2. kubectl get pod cpu-demo-2 --namespace(cpu-example)
3. kubectl describe pod cpu-demo-2 --namespace(cpu-example)
4. kubectl delete pod cpu-demo-2 --namespace(cpu-example)

CleanUp Scenario

- kubectl delete namespace cpu-example
-
-

Configure Quality of Service for Pods

- 1.QoS classes:
 - 1.Guaranteed
 - 2.Burstable
 - 3.BestEffort

1.kubectl create namespace qos-example

Create a Pod that gets assigned a QoS class of Guaranteed

1. For a Pod to be given a QoS class of Guaranteed:
 1. Every Container in the Pod must have a memory limit and a memory request, and they must be the same.
 2. Every Container in the Pod must have a CPU limit and a CPU request, and they must be the same.

Assign CPU Resources to Containers and Pods

Senario-6

1. kubectl create -f https://k8s.io/examples/pods/qos/qos-pod.yaml --namespace=qos-example
2. kubectl get pod qos-demo --namespace=qos-example –output=yaml
3. kubectl delete pod qos-demo --namespace=qos-example

Create a Pod that gets assigned a QoS class of Burstable

1. A Pod is given a QoS class of Burstable if:
 1. The Pod does not meet the criteria for QoS class Guaranteed.
 2. At least one Container in the Pod has a memory or CPU request.

Create a Pod that gets assigned a QoS class of Burstable

Senario-7

1. `kubectl create -f https://k8s.io/examples/pods/qos/qos-pod-2.yaml --namespace=qos-example`

2. `kubectl create -f https://k8s.io/examples/pods/qos/qos-pod-2.yaml --namespace=qos-example`

3. `kubectl delete pod qos-demo-2 --namespace=qos-example`

Create a Pod that gets assigned a QoS class of BestEffort

1. A Pod is given a QoS class of BestEffort if:
 1. For a Pod to be given a QoS class of BestEffort, the Containers in the Pod must not have any memory or CPU limits or requests.

Create a Pod that gets assigned a QoS class of BestEffort

Senario-8

1. `kubectl create -f https://k8s.io/examples/pods/qos/qos-pod-3.yaml --namespace=qos-example`

2. `kubectl get pod qos-demo-3 --namespace=qos-example --output=yaml`

3. `kubectl delete pod qos-demo-3 --namespace=qos-example`

Create a Pod that has two Containers

Senario-9

1. `kubectl create -f https://k8s.io/examples/pods/qos/qos-pod-4.yaml --namespace=qos-example`

2. `kubectl get pod qos-demo-4 --namespace=qos-example --output=yaml`

3. `kubectl delete pod qos-demo-4 --namespace=qos-example`

CleanUp Scenario

- kubectl delete namespace qos-example

Configure a Pod to Use a Volume for Storage

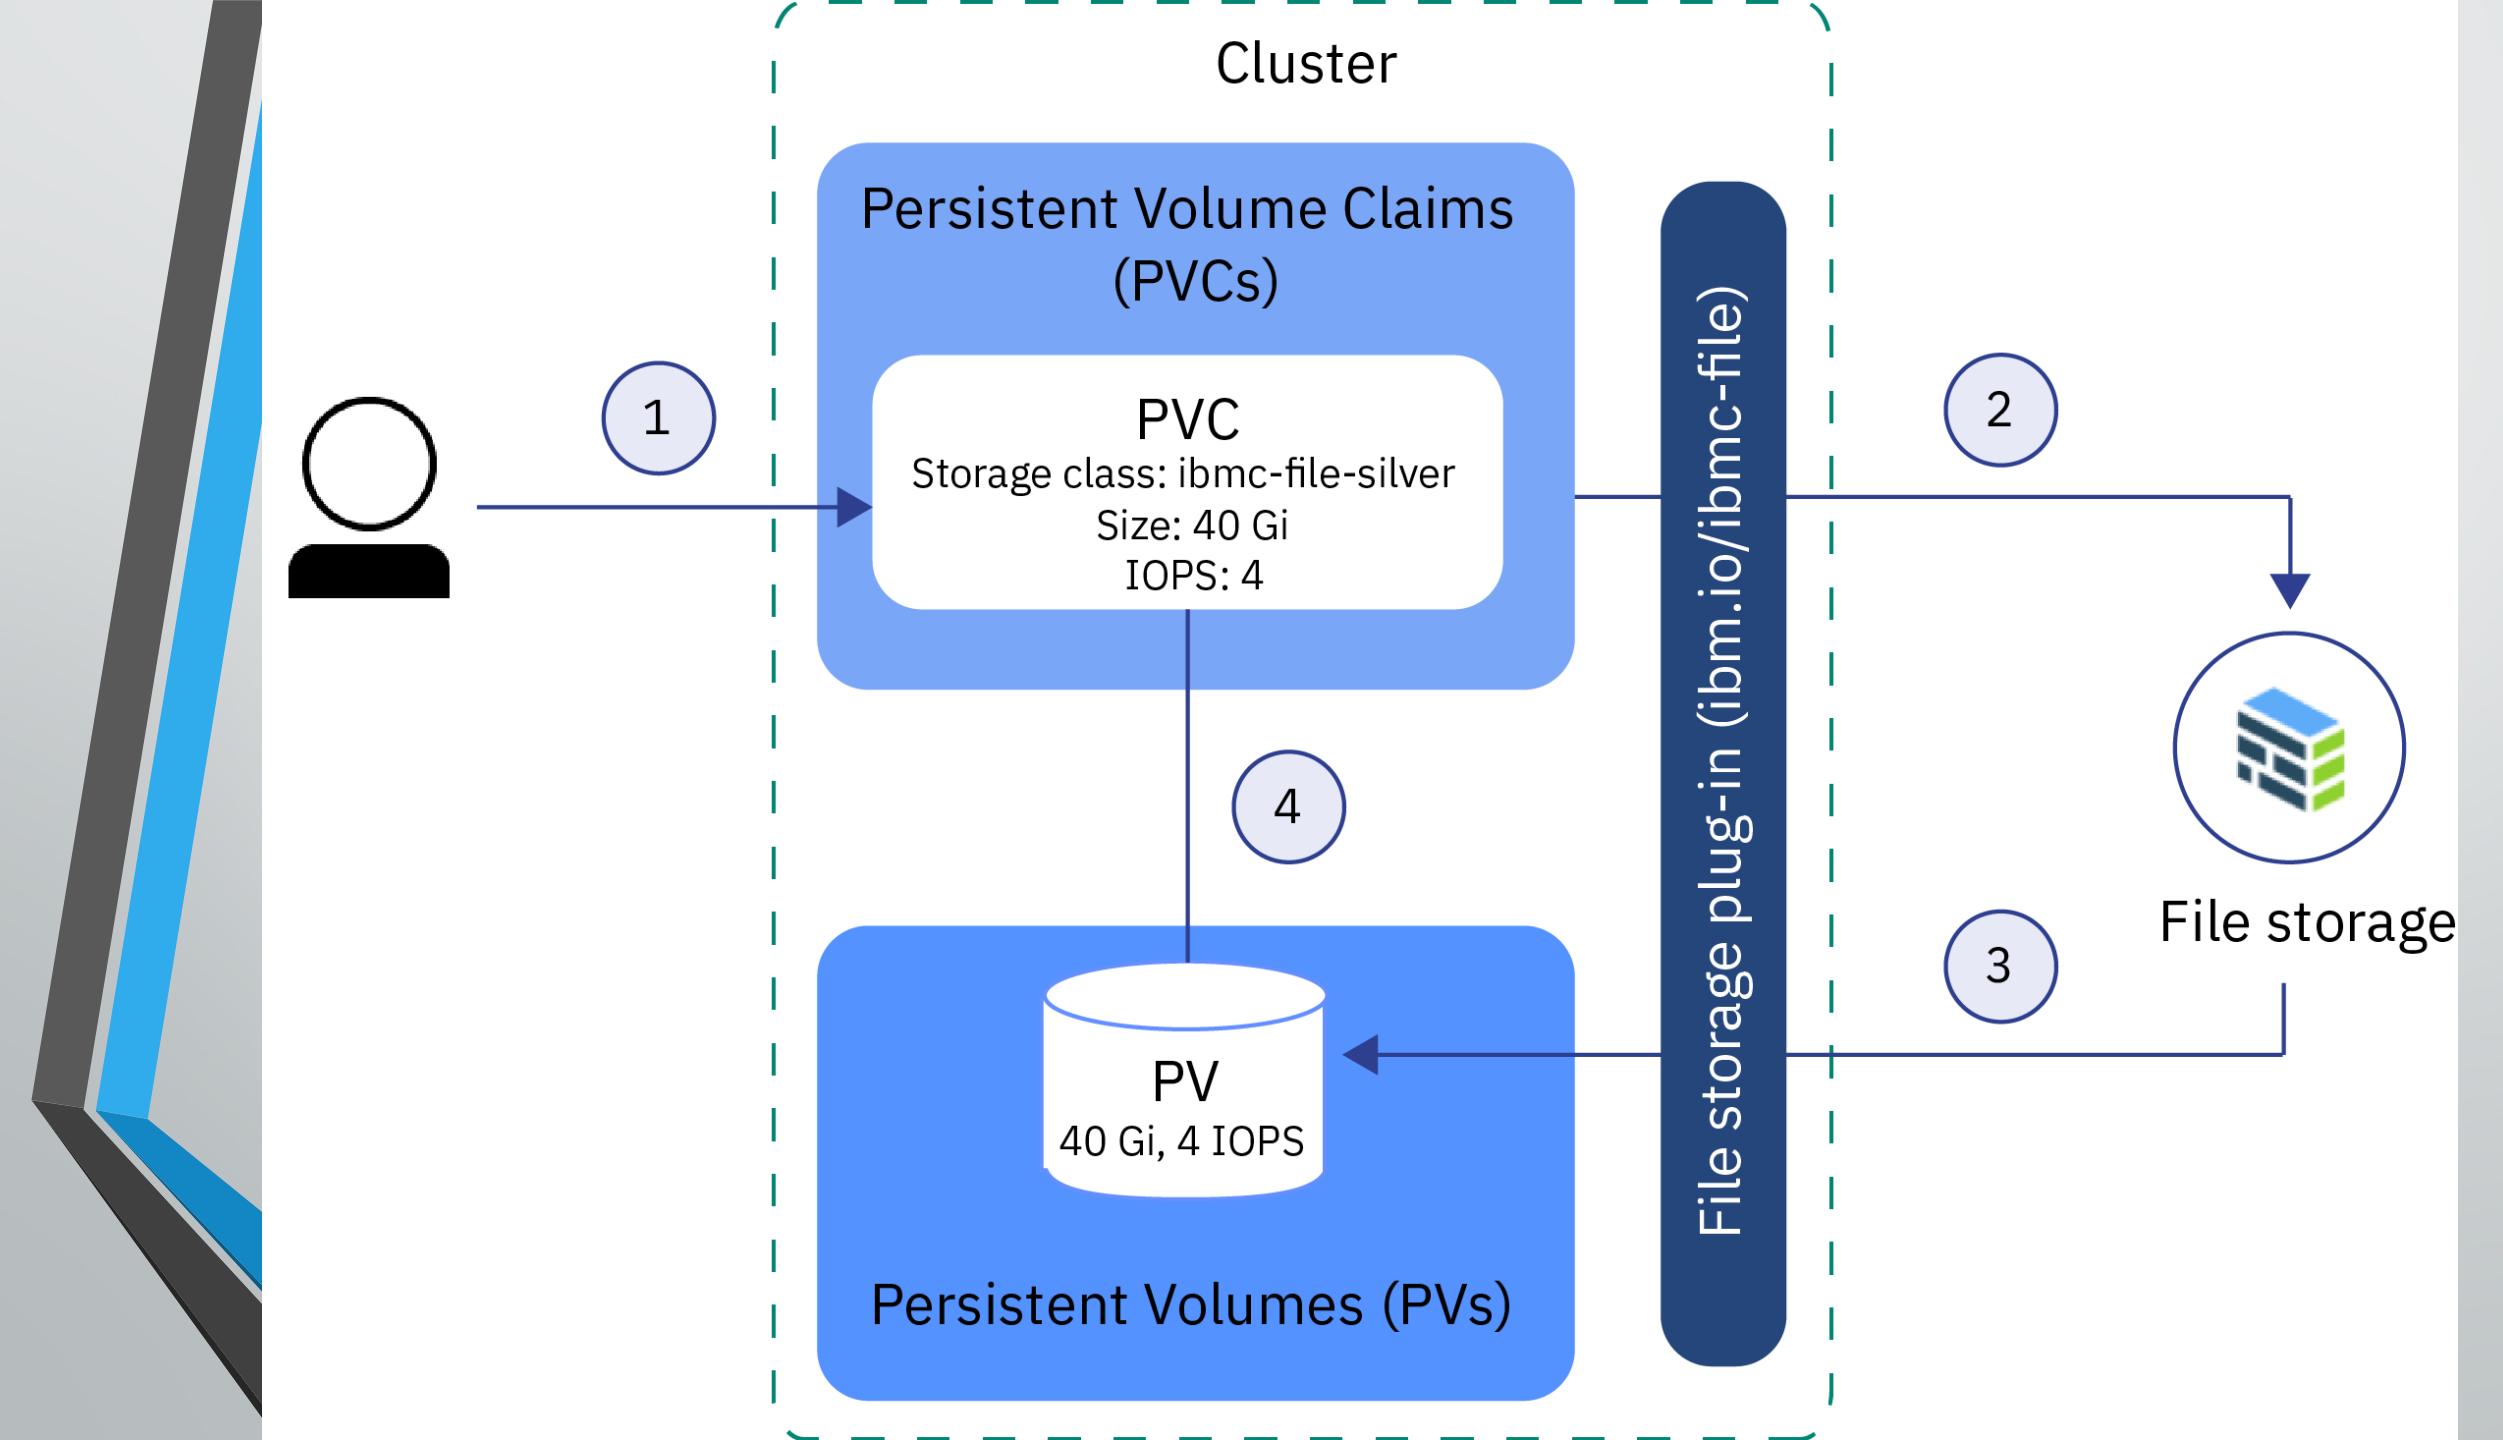
Configure a volume for a Pod

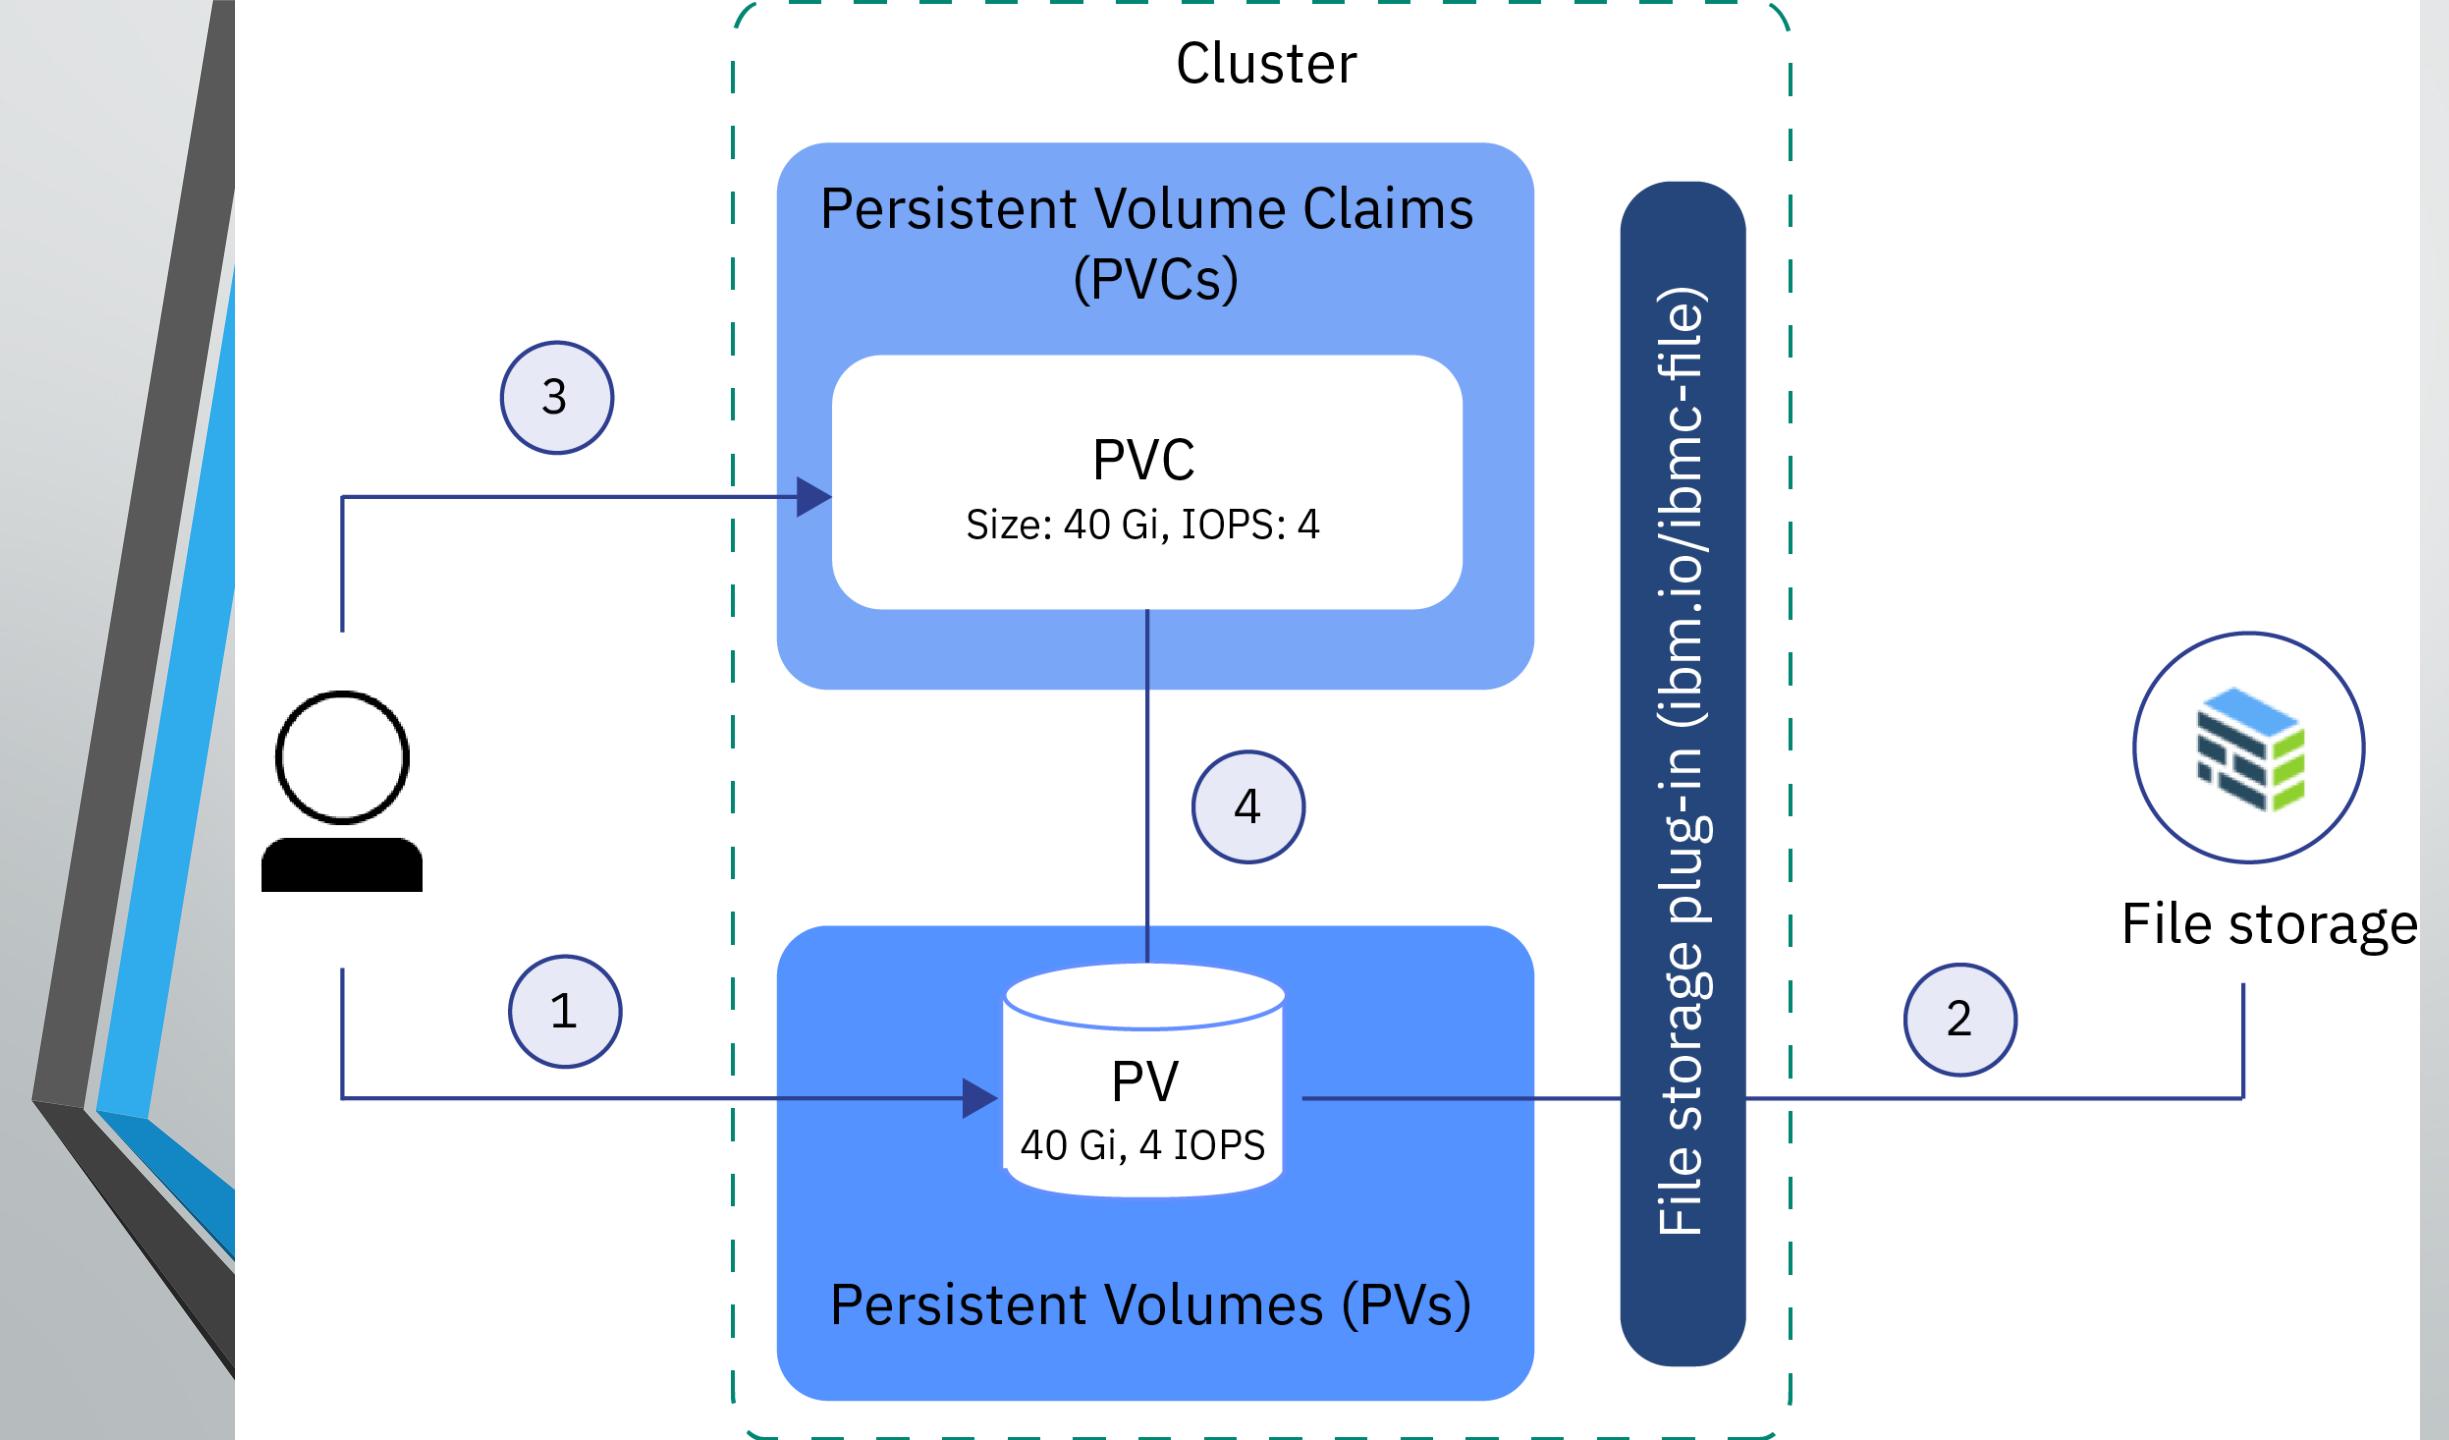
Senario-10

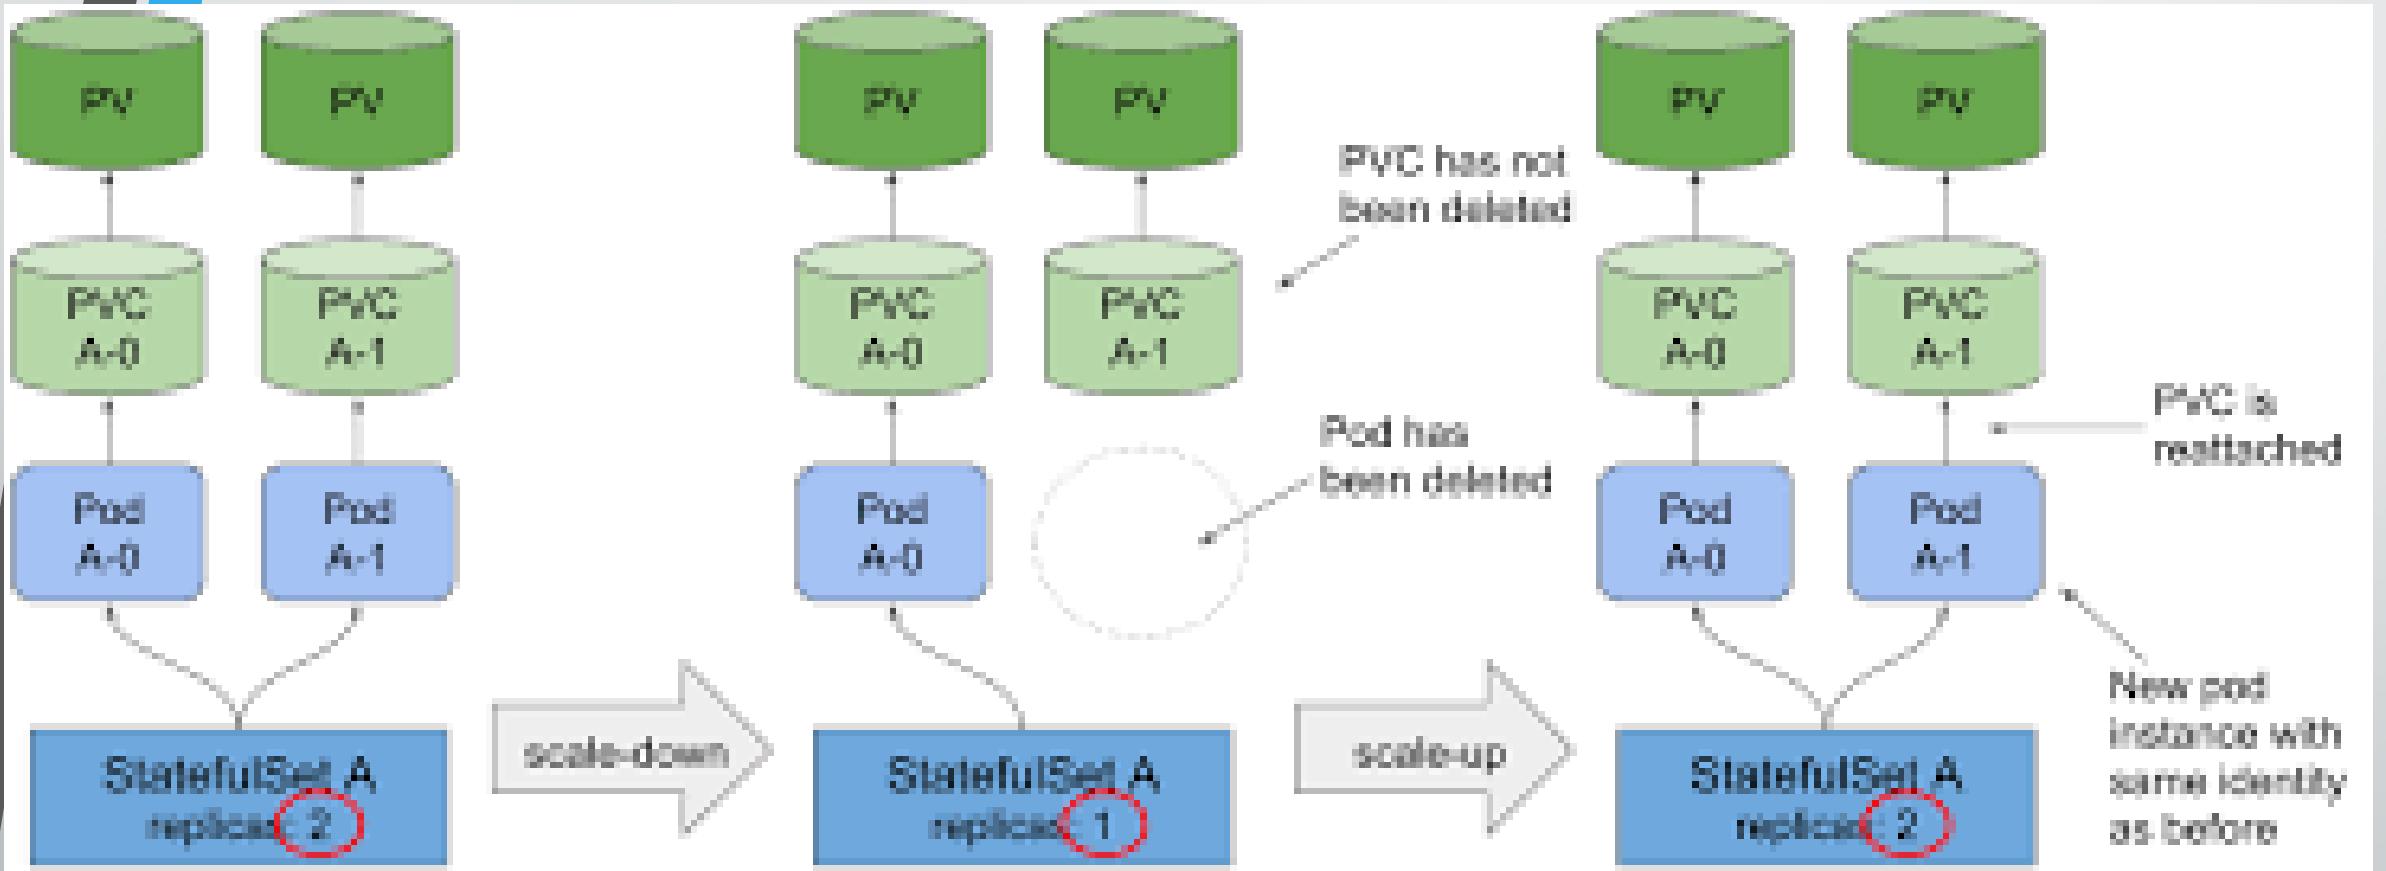
1. `kubectl create -f https://k8s.io/examples/pods/storage/redis.yaml`
2. `kubectl get pod redis --watch`
- 3.
4. `kubectl exec -it redis -- /bin/bash`
5. `kubectl delete pod redis`



Configure a Pod to Use a PersistentVolume for Storage







Create a PersistentVolume

Senario-11

1. mkdir /mnt/data
2. echo 'Hello from Kubernetes storage' > /mnt/data/index.html
3. kubectl create -f <https://k8s.io/examples/pods/storage/pv-volume.yaml>
4. kubectl get pv task-pv-volume

Access Modes

1. `ReadWriteOnce`: The Volume can be mounted as read-write by a single node.
2. `ReadOnlyMany`: The Volume can be mounted read-only by many nodes.
3. `ReadWriteMany`: The Volume can be mounted as read-write by many nodes. PersistentVolumes that are backed by Compute Engine persistent disks don't support this access mode.
- 4.

Create a PersistentVolumeClaim

Senario-12

1. Step-1:
 1. mkdir /mnt/data
 2. echo 'Hello from Kubernetes storage' > /mnt/data/index.html
 3. kubectl create -f <https://k8s.io/examples/pods/storage/pv-volume.yaml>
 4. kubectl get pv task-pv-volume
2. Step-2:
 1. kubectl create -f <https://k8s.io/examples/pods/storage/pv-claim.yaml>
 2. kubectl get pv task-pv-volume
 3. kubectl get pvc task-pv-claim
3. Step-3
 1. kubectl create -f <https://k8s.io/examples/pods/storage/pv-pod.yaml>
 2. kubectl get pod task-pv-pod
 3. kubectl exec -it task-pv-pod -- /bin/bash
 4. root@task-pv-pod:# apt-get update
 5. root@task-pv-pod:# apt-get install curl
 6. root@task-pv-pod:# curl localhost

Configure Default Memory Requests and Limits for a Namespace

Configure Default Memory Requests and Limits for a Namespace

Senario-13

1. Create a namespace
 1. `kubectl create namespace default-mem-example`
2. Create a LimitRange and a Pod
 1. `kubectl create -f https://k8s.io/examples/admin/resource/memory-defaults.yaml --namespace=default-mem-example`
 2. `kubectl create -f https://k8s.io/examples/admin/resource/memory-defaults-pod.yaml --namespace=default-mem-example`
3. `kubectl get pod default-mem-demo --output=yaml --namespace=default-mem-example`
4. `kubectl delete pod default-mem-demo --namespace=default-mem-example`

What if you specify a Container's limit, but not its request?

Senario-14

1. `kubectl create -f https://k8s.io/examples/admin/resource/memory-defaults-pod-2.yaml --namespace=default-mem-example`
2. `kubectl get pod default-mem-demo-2 --output=yaml --namespace=default-mem-example`
- 3.

What if you specify a Container's request, but not its limit?

Senario-15

1. kubectl create -f <https://k8s.io/examples/admin/resource/memory-defaults-pod-3.yaml> --namespace=default-mem-example
- 2.
3. kubectl get pod default-mem-demo-3 --output=yaml --namespace=default-mem-example
- 4.
- 5.

Configure a Pod Quota for a Namespace

Configure a Pod Quota for a Namespace

Senario-16

1. Create a namespace
 1. `kubectl create namespace quota-pod-example`
2. Create a ResourceQuota
 1. `kubectl create -f https://k8s.io/examples/admin/resource/quota-pod.yaml --namespace=quota-pod-example`
 2. `kubectl get resourcequota pod-demo --namespace=quota-pod-example --output=yaml`
3. `kubectl create -f https://k8s.io/examples/admin/resource/quota-pod-deployment.yaml --namespace=quota-pod-example`
4. `kubectl get deployment pod-quota-demo --namespace=quota-pod-example --output=yaml`
5. `kubectl delete namespace quota-pod-example`
- 6.



Assign Pods to Nodes

Configure a Pod Quota for a Namespace

Senario-17

1. kubectl get nodes
2. kubectl label nodes <your-node-name> disktype=ssd
3. kubectl get nodes --show-labels
- 4.
5. kubectl create -f <https://k8s.io/examples/pods/pod-nginx.yaml>
6. kubectl get pods --output=wide

Running Automated Tasks with a CronJob

Creating a Cron Job

Senario-18

1. kubectl create -f <https://raw.githubusercontent.com/kubernetes/website/master/content/en/examples/application/job/cronjob.yaml>
- 2.
3. kubectl get cronjob hello
4. kubectl get jobs --watch